



This is essentially describing **LLM/RAG response caching** and **prompt caching**, which solve related but different problems.

1. RAG response caching

The flow is:




For example:





Next time:



So you avoid:

embedding/retrieval → prompt construction → LLM generation

and get a much faster response.

2. The important distinction: RAG cache vs Prompt Cache

These are often confused.

Feature	RAG Response Cache	Prompt Caching
What is cached?	Final/generated answer	Reusable input tokens
Cache location	Redis/SQLite/etc.	LLM provider infrastructure
LLM called on HIT?	Usually No	Yes
Retrieval required?	No	Depends
Reduces latency	✓ Significant	✓
Reduces LLM cost	✓	✓ Significant

Feature	RAG Response Cache	Prompt Caching
Handles slightly different questions	Usually no	✓
Best for	Repeated questions	Large repeated context/system prompts

RAG response cache

```

Question
↓
Cache
↓ HIT
Final Answer

```



Anthropic-style prompt caching

```

Large Document / System Prompt
↓
Cached Prefix
↓
User Question → LLM
↓
Answer


```




The model still runs for every question.

3. For your Advanced RAG architecture

Since you're working with **FastAPI** + **Qdrant** + **Ollama**, I'd recommend a more production-oriented design than simply caching the question.





4. But there's an even better approach: Semantic Cache

Exact caching has a limitation.

These two questions are different strings:


"What is Snappy?"

"What does Snappy do?"





An exact cache treats them as different.

A **semantic cache** can determine that they are asking essentially the same thing.



So your architecture can become:





This is particularly useful for your RAG system.

5. Cache invalidation is the critical part

The lecture correctly highlights this.

Imagine your documentation changes:

Documentation v1
↓
"What is Product X?"
↓
Cached answer:
"Product X supports A, B and C"



Then documentation changes:

Documentation v2
↓
Product X no longer supports C



But Redis still contains:

"What is Product X?"
↓
"A, B and C" ❌



That's dangerous for RAG.

A good solution is **versioned cache keys**:

rag:v1:website123:question_hash



After updating documents:

rag:v2:website123:question_hash



Now old answers are automatically bypassed.

You can also use:

```
TTL
+
Document Version
+
Model Version
+
Prompt Version
```



For example:


```
cache_key =
    hash(
        tenant_id +
        document_version +
        prompt_version +
        model_version +
        normalized_question
    )
```



6. Don't cache every answer blindly

This is especially important for your **RAG guardrails** architecture.

A safer pipeline is:



You can apply conditions such as:

```
faithfulness > threshold
context_precision > threshold
answer_relevancy > threshold
injection_detected == false
blocked_reason == null
```



Only then:


```
Redis SET cache_key answer
```



This prevents a hallucinated answer from becoming a **permanent cached hallucination**.

7. Where LangChain fits

Conceptually:



LLM caching primarily avoids repeated LLM calls for identical inputs.

But for your production RAG system, I'd separate:

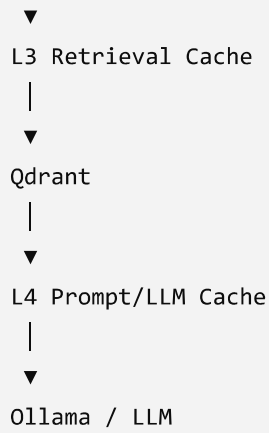
```
L1 → Exact Response Cache
L2 → Semantic Cache
L3 → Retrieval Cache
L4 → LLM/Prompt Cache
```



For example:

```
User
|
▼
L1 Exact Cache — HIT → Answer
| MISS
▼
L2 Semantic Cache — HIT → Answer
| MISS
```





That gives you a much stronger **production-grade RAG caching architecture** than simply putting `InMemoryCache` in front of the LLM.

The key takeaway is:

Response caching caches the answer; semantic caching caches the meaning; prompt caching caches reusable input context.

For your RAG platform, **semantic response caching + cache versioning + guardrail validation + Redis** would be the most interesting implementation to demonstrate in an architecture/interview project.



ChatGPT can make mistakes. Check important info.