

How Calculates the Similarity Between This Query Vector and Document Vectors

Perfect 🌟 Let's explain this in **very simple way**, no heavy math.

You saw this flow:

Query → Query Vector → Compare with Document Vectors → Top-K Results

Now let's understand **how similarity is calculated** in easy language.

Step 1: Convert Text into Numbers (Embedding)

Example:

User searches:

"red shoes"

Embedding model converts it into something like:

Query Vector (Q)

[0.6, 0.3, 0.1]

Each product in database is already converted into numbers:

Doc1 (red sneakers) → [0.58, 0.29, 0.12]

Doc2 (blue sandals) → [0.10, 0.80, 0.05]

Doc3 (running shoes) → [0.55, 0.25, 0.15]

These numbers represent **meaning**, not words.

● Step 2: Think of Each Vector as a Point in Space

Imagine 3D space (just for understanding):

- Doc2
- Doc1 ● Doc3
- Query

If two points are close → similar meaning

If far → different meaning

The database just checks:

“Which points are closest to the Query point?”

That's it.

● Step 3: How Does It Check “Closeness”?

Instead of complicated math thinking, imagine:

⌚ Method 1: Angle Between Arrows (Cosine Similarity)

Think of vectors like arrows starting from the same center:

Doc1 ↗
/
/
•
/
↗ Query

If arrows point in same direction → very similar

If arrows go opposite → very different

Database checks:

👉 “Are these arrows pointing in similar direction?”

More same direction = higher similarity score (like 0.95)

Method 2: Distance Between Points

Imagine measuring with a ruler:

Query •----• Doc1 (short distance → similar)

Query •-----• Doc2 (long distance → not similar)

Short distance = high similarity

Long distance = low similarity

Step 4: Ranking (Top-K)

After calculating similarity score with all documents:

Doc1 → 0.95

Doc3 → 0.92

Doc2 → 0.30

If K = 2

Database returns:

1. Doc1

2. Doc3

That's called **Top-K Retrieval**.

What Actually Happens Internally (Simple Version)

1. Query text → converted to numbers
2. For each document:
 - Compare direction OR distance
3. Give each document a similarity score
4. Sort by highest score
5. Return top K

No keyword matching.

No SQL WHERE.

Just geometry.

Very Simple Real-Life Analogy

Imagine:

You are in a room with 100 people.

You ask:

"Who likes red shoes?"

Instead of asking each person verbally, you:

- Measure personality similarity
- Stand near people with similar interests
- Closest people = most similar

Vector DB does exactly that — but in high-dimensional math space.

Why It Feels Smart

Because embedding models place similar meanings close together:

"red shoes"

"red sneakers"

"buy running shoes"

All cluster together.

But:

"blue hat"

"cooking recipe"

Are far away.

 **Important**

The database is NOT understanding language.

It is only:

Comparing numeric patterns in multi-dimensional space.

But because embeddings capture meaning, it *looks* intelligent.