

What are Embeddings?



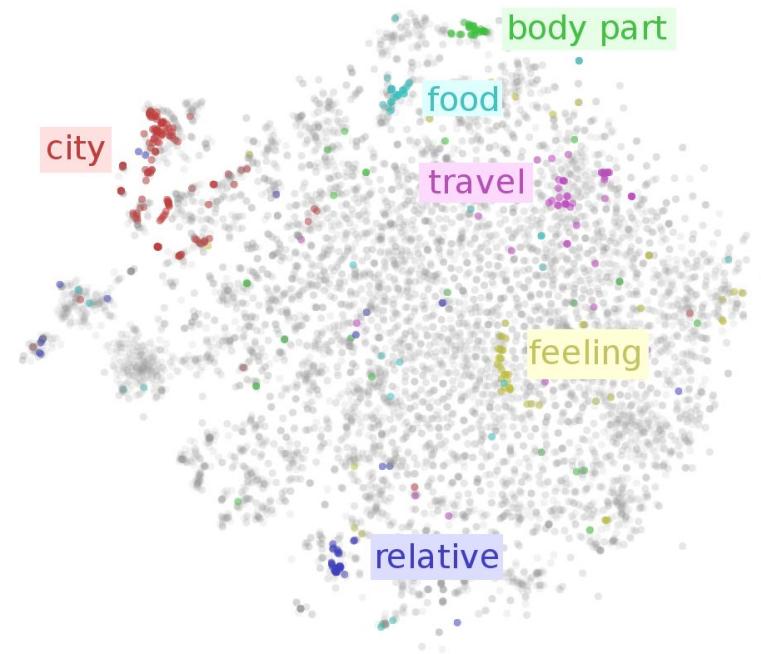
Learn word embeddings, how to use them and their various use cases

What are Embeddings?

Embeddings are numerical representations of data (text, images, audio) in a **high-dimensional vector space**.

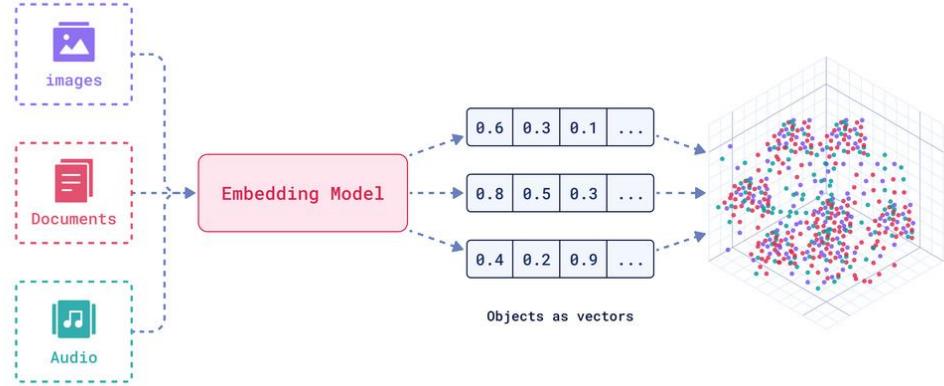
They capture **semantic meaning**, making **similar concepts appear close together** mathematically.

Core technology enabling many AI applications, especially in natural language processing.

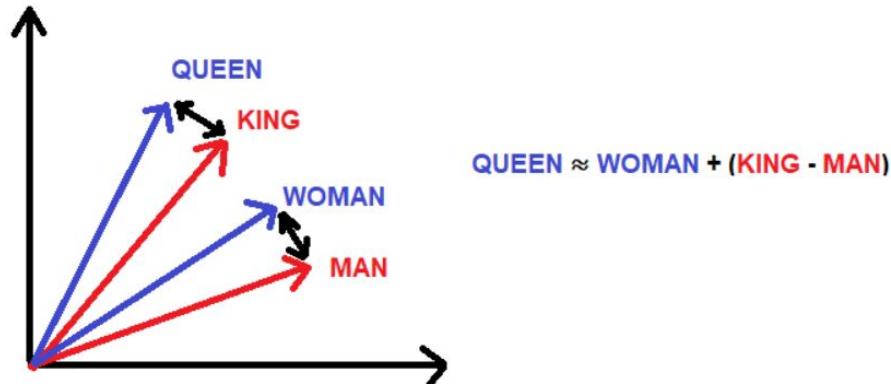


How Embeddings Work

1. **Tokenization:** Break text into tokens (words, subwords)
2. **Vector Creation:** Convert tokens/images into high-dimensional vectors
3. **Contextual Understanding:** Position vectors based on meaning in context



How Embeddings Work



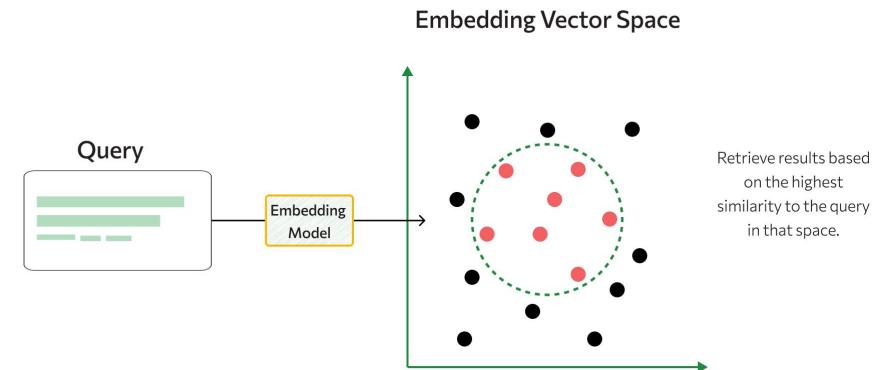
As embeddings represent a **space within multiple dimensions**, different embeddings will either **be closer or further away from other embeddings**.

Common Embedding Models

- [OpenAI:](#) text-embedding-ada-002,
text-embedding-3-small/large
- [Cohere:](#) embed-english-v3.0,
embed-multilingual
- [Sentence Transformers:](#)
all-MiniLM-L6-v2, all-mpnet-base-v2
- [Google:](#) PaLM, BERT, Universal Sentence
Encoder

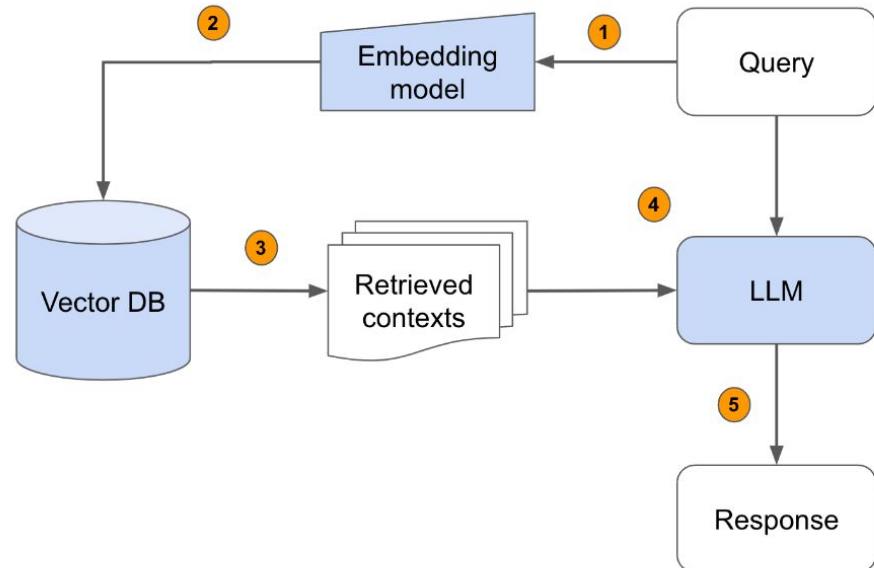
Applications in Prompt Engineering - Semantic Search

- Match queries to documents based on meaning, not just keywords
- Enable more intuitive information retrieval



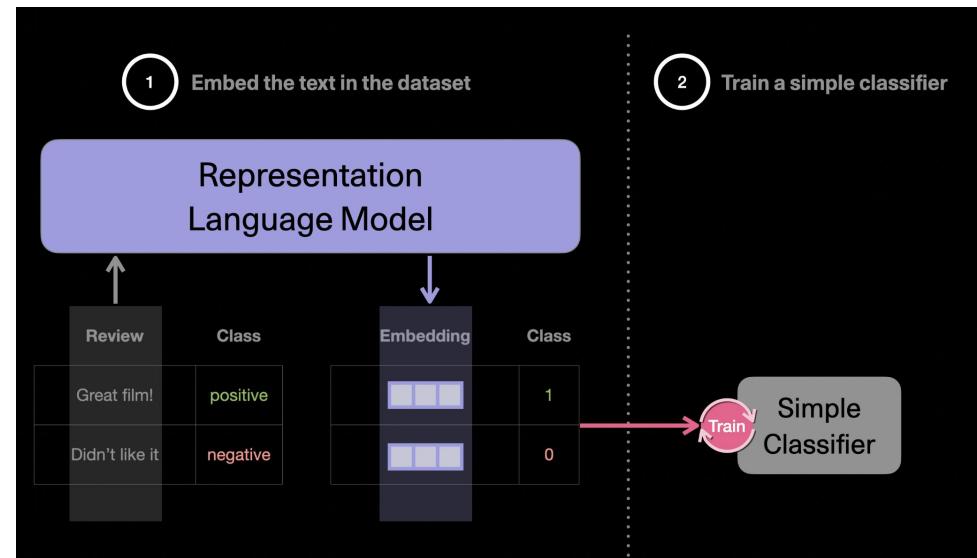
Applications - Retrieval-Augmented Generation (RAG)

- Find relevant information from a knowledge base
- Incorporate into prompts/chat history to provide context for LLMs



Applications in Prompt Engineering - Text Classification

- Categorize content based on semantic patterns
- Power content recommendation systems



Embeddings in Practice

```
from openai import OpenAI
client = OpenAI()

response = client.embeddings.create(
    input="Your text string goes here",
    model="text-embedding-3-small"
)

print(response.data[0].embedding)
```

```
{
  "object": "list",
  "data": [
    {
      "object": "embedding",
      "index": 0,
      "embedding": [
        -0.006929283495992422,
        -0.005336422007530928,
        -4.547132266452536e-05,
        -0.024047505110502243
      ],
    }
  ],
  "model": "text-embedding-3-small",
  "usage": {
    "prompt_tokens": 5,
    "total_tokens": 5
  }
}
```

Measuring Similarity

- **Cosine similarity:** measures angle between vectors
- **Euclidean distance:** measures straight-line distance
- **Dot product:** simple multiplication of vectors

Cosine Similarity - The Heart of Embedding Comparisons

- Mathematical measure of similarity between two vectors regardless of their magnitude
- Calculates the cosine of the angle between vectors in a multi-dimensional space
- Ranges from -1 (opposite direction) to 1 (same direction), with 0 indicating orthogonality

Cosine Distance:

$$d_{cos}(x,y) = 1 - \frac{\sum_{i=0}^{n-1} x_i y_i}{\sqrt{\sum_{i=0}^{n-1} x_i^2} \times \sqrt{\sum_{i=0}^{n-1} y_i^2}}$$

Cosine Similarity - Coding Example

```
import numpy as np
from sklearn.metrics.pairwise import cosine_similarity

# Two embedding vectors
embedding1 = np.array([0.2, 0.5, 0.3, 0.8, 0.1])
embedding2 = np.array([0.3, 0.4, 0.2, 0.7, 0.2])

# Calculate similarity
similarity = cosine_similarity([embedding1], [embedding2])[0][0]
print(f"Similarity score: {similarity:.4f}") # Output: Similarity score: 0.9819
```

Advanced Techniques & Future Directions

- **Domain Specialization:** Fine-tuned embeddings for industry-specific language and applications
- **Multimodal Integration:** Models like CLIP and ImageBind creating unified vector spaces across text, images, and audio
- **Hybrid Approaches:** Combining neural embeddings with traditional methods (BM25/TF-IDF) for optimal results
- **Key Challenge:** Keeping embeddings updated with new knowledge without complete retraining

Next Steps



Let's get hands on and practice learning how to create embeddings within OpenAI, and using cosine similarity to compare different embeddings!